

# Image Inpainting

Dwaipayan Haldar  
BEE-IV, 002110801187  
Dept. of Electrical Engineering  
Jadavpur University

4th April 2025

# Outline

- 1 Introduction
- 2 Classification
- 3 Exemplar based Texture Synthesis
- 4 General Adversarial Networks
- 5 Conclusion
- 6 References

# Introduction

# What is Image Inpainting?

Image inpainting originated from an ancient technique performed by artists to restore damaged paintings or photographs with small defects such as scratches, cracks, dust and spots to maintain its quality to as close to the original as possible.



Figure 1: Hand inpainting performed by an artist. Image courtesy of Thottam (2015).

# Object Removal

Image inpainting removes unwanted objects by filling the gaps with realistic textures using deep learning (GANs, diffusion models) or traditional methods, ensuring seamless restoration.



(a) Real image



(b) Modified image

**Figure 2:** An example of a people removal method where the person has been manually selected in the real image (a), and then automatically removed in the second image (b) by filling the region concerning the person using an exemplar-based image inpainting method. Reprinted from Criminisi et al. (2004).

# Example Problem



Figure 3: Photo of Dwaipayan Haldar

# Example Problem



Figure 4: Photo of Dwaipayan Haldar distorted by Manas Sarkar

# Example Problem



Figure 5: Ghibli of Dwaipayan Haldar

# Example Problem



Figure 6: Ghibli of Dwaipayan Haldar distorted by Pramit Khamrui

# Classification

# Classification of different Inpainting Techniques

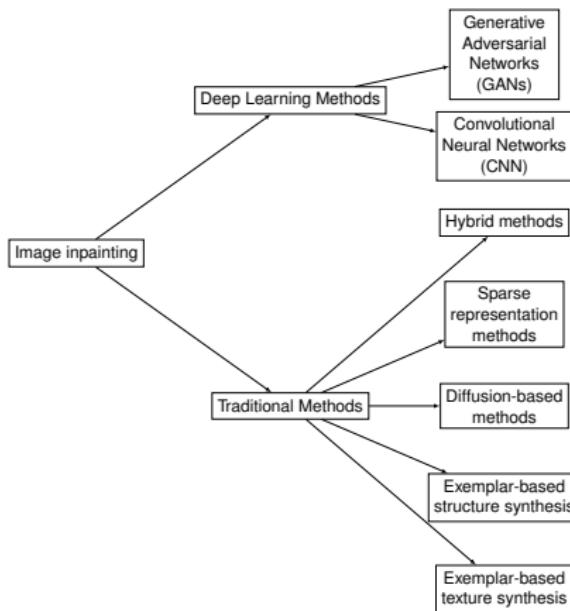


Figure 7: Hierarchical representation of image inpainting techniques in two main categories

# Classification of different Inpainting Techniques

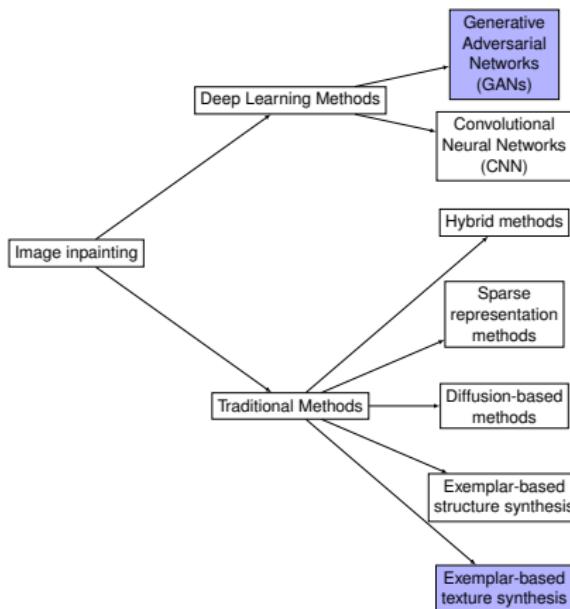


Figure 8: Hierarchical representation of image inpainting techniques in two main categories

# Exemplar Based Texture Synthesis

# Definitions

## What is a texture?

A texture is a visual pattern on an infinite 2-D plane with a stationary distribution at some scale (Efros and Leung, 1999). This pattern refers to the feel (smooth, rough) of the image surface.

## Types of textures

Textures have been traditionally classified as either regular (consisting of repeated texels) or stochastic (without explicit texels). However, almost all real-world textures lie somewhere in between these two extremes and should be captured with a single model.

## Assumptions

- **Markov Random Field (MRF) Model:** The algorithm assumes that texture can be modeled as a Markov Random Field. This means that the probability distribution of a pixel's brightness value, given the brightness values of its spatial neighborhood, is independent of the rest of the image. In simpler terms, a pixel's value depends only on its neighbors.
- **Stationary Texture:** The usual assumption, which this paper also makes, is that the sample texture is large enough to capture the stationarity of the texture. "Stationarity" in this context means that the statistical properties of the texture are consistent across the image.

# Overview of Algorithm

- Let  $I$  be an image being synthesized from a texture sample image  $I_{smp}$ , where  $I_{real}$  is the real infinite texture.
- Let  $p \in I$  be a pixel, and  $\omega(p) \subset I$  be a square image patch of width  $w$  centered at  $p$ .
- Let  $d(\omega_1, \omega_2)$  denote some perceptual distance between two patches.
- Assume all pixels in  $I$  except  $p$  are known.
- Based on the MRF model, assume that  $p$  is independent of  $I \setminus \omega(p)$  given  $\omega(p)$ .

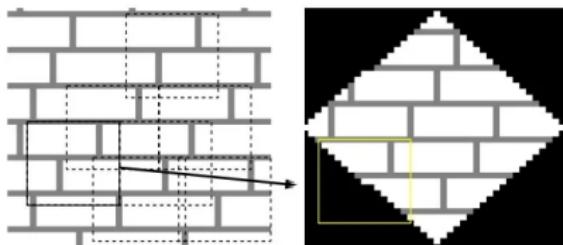


Figure 9: Algorithm Overview

# Overview of the algorithm

- Define the set:

$$S(p) = \{\omega' \subset I_{real} : d(\omega', \omega(p)) = 0\}$$

containing all occurrences of  $\omega(p)$  in  $I_{real}$ .

- Estimate the conditional probability density function (pdf) of  $p$  using a histogram of all center pixel values in  $S(p)$ .
- Since only a finite sample  $I_{smp}$  is available, there may not be exact matches for  $\omega(p)$ .
- Use a heuristic to find a plausible subset  $S_0(p) \subset S(p)$  for sampling.
- Implement a nearest neighbor technique: find the best match  $\omega_{best} = \arg \min_{\omega} d(\omega(p), \omega)$  in  $I_{smp}$ .

# Overview of the algorithm

- Include all patches  $\omega$  satisfying:

$$d(\omega(p), \omega) < (1 + \epsilon)d(\omega(p), \omega_{best})$$

where  $\epsilon = 0.1$ .

- Use the center pixel values of patches in  $S_0(p)$  to construct a histogram for  $p$ .
- Sample from this histogram, either uniformly or weighted by  $d$ .
- Define a suitable distance metric  $d$ , such as normalized sum of squared differences (SSD) with Gaussian kernel.

$$d_{ssd} \cdot G = \frac{\sum_{i,j} G(i,j) \cdot (F_{i,j} - I_{i,j})^2}{\sqrt{\sum_{i,j} F_{i,j}^2 \cdot \sum_{i,j} I_{i,j}^2}}$$

# Hole filling with Texture Synthesis

123	89	150	130	110
80	90	70	40	100
77	67	59	130	120
89	30	78	90	100
20	50	70	80	120

123	89	150	130	110	130
80			40	100	120
77			130	120	100
89	30	78	90	100	70
20	50	70	80	120	50
110	90	89	76	85	100

# Hole filling with Texture Synthesis

123	89	150	130	110
80	90	70	40	100
77	67	59	130	120
89	30	78	90	100
20	50	70	80	120

123	89	150	130	110	130
80			40	100	120
77			130	120	100
89	30	78	90	100	70
20	50	70	80	120	50
110	90	89	76	85	100

# Hole filling with Texture Synthesis

123	89	150	130	110
80	90	70	40	100
77	67	59	130	120
89	30	78	90	100
20	50	70	80	120

123	89	150	130	110	130
80			40	100	120
77			130	120	100
89	30	78	90	100	70
20	50	70	80	120	50
110	90	89	76	85	100

# Hole filling with Texture Synthesis

123	89	150	130	110
80	90	70	40	100
77	67	59	130	120
89	30	78	90	100
20	50	70	80	120

123	89	150	130	110	130
80			40	100	120
77			130	120	100
89	30	78	90	100	70
20	50	70	80	120	50
110	90	89	76	85	100

# Results



Figure 10: Inpainted Image with 5\*5 window



Figure 11: Inpainted Image with 13\*13 window



Figure 12: Inpainted Image with 25\*25 window

# Results



Figure 13: Original Image



Figure 14: Inpainted Image with 25\*25 window

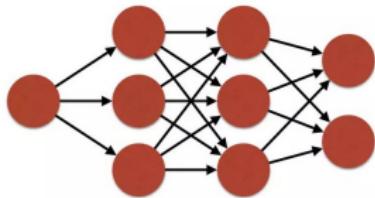
## Some Key Points

- Very Simple
- Surprisingly good results
- But very slow
- Limited in terms of mask size, accuracy and sometimes efficiency

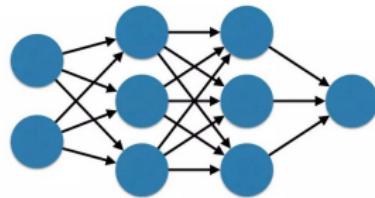
# General Adversarial Networks

# Introduction

GAN consists of two neural networks. One is called the generator, and the other is the discriminator.



Generator

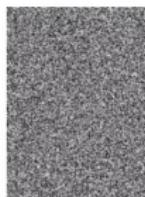


Discriminator

# Introduction



Generator

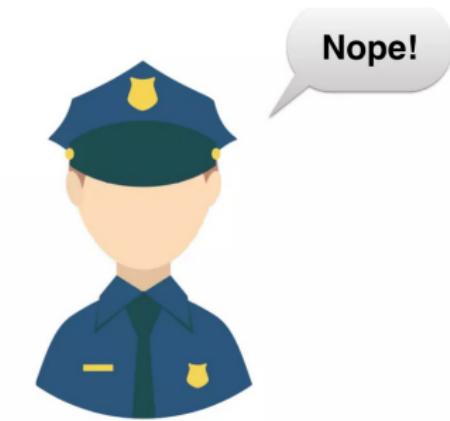


Discriminator

# Introduction



Generator



Discriminator

# Introduction

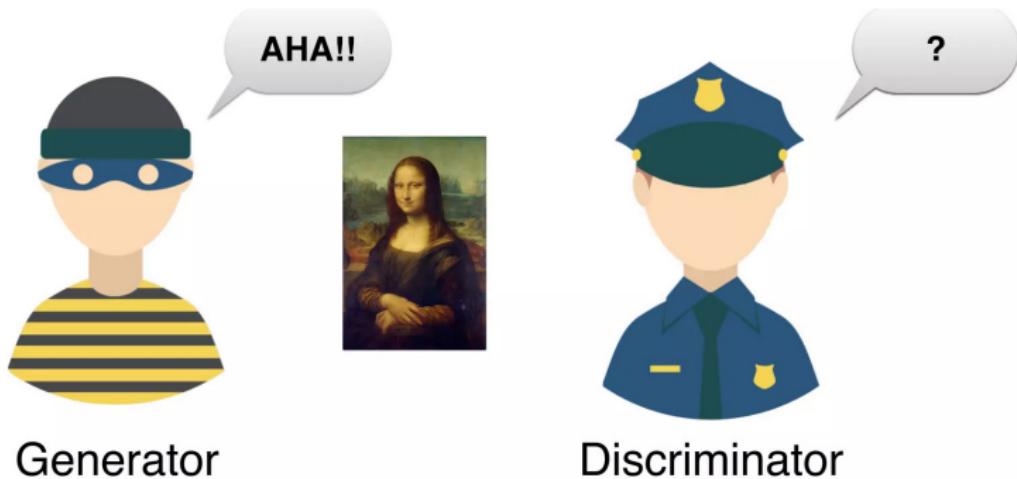


Generator

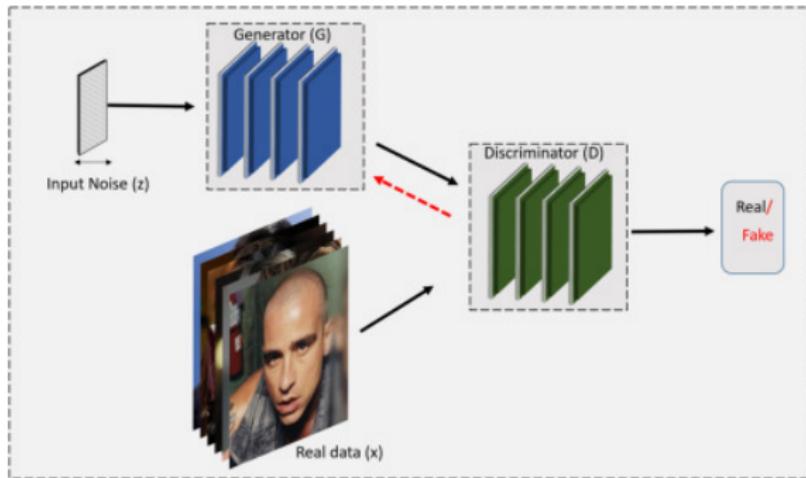


Discriminator

# Introduction

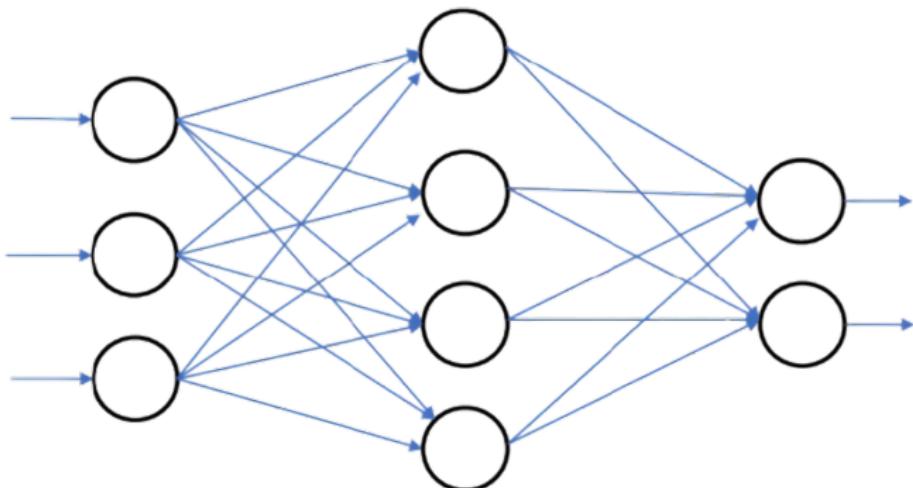


# GAN Block Architecture

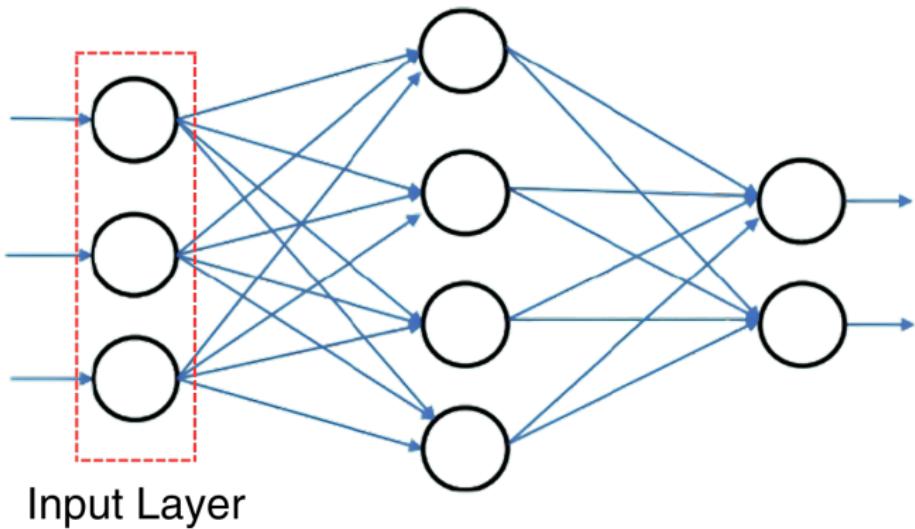


**Figure 15:** An example GAN block. The input of the generator ( $z$ ) is sampled from a random noise vector and the input of the discriminator ( $x$ ) is sampled from real data distribution.

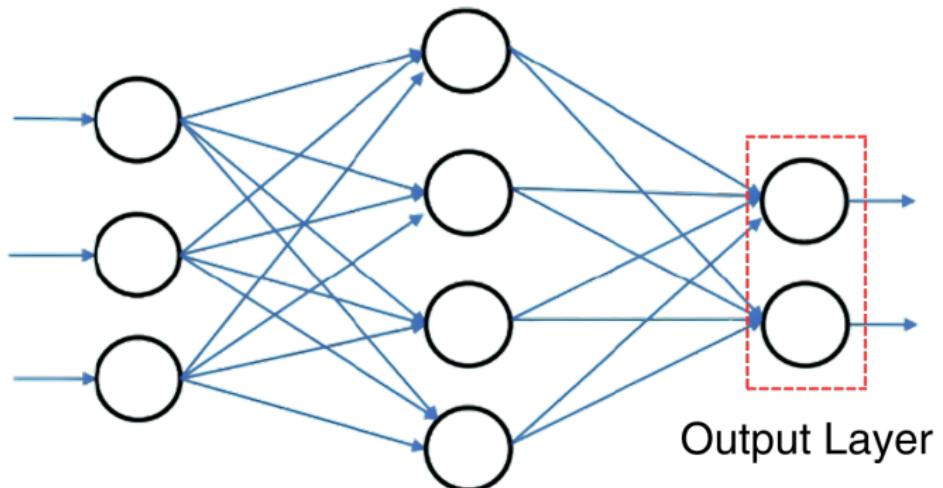
# Neural Networks



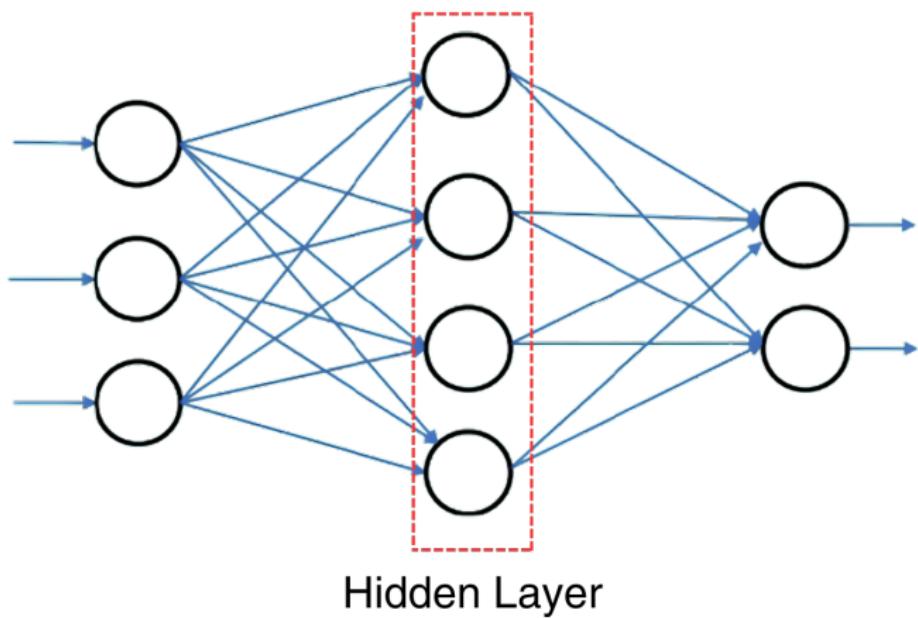
# Neural Networks



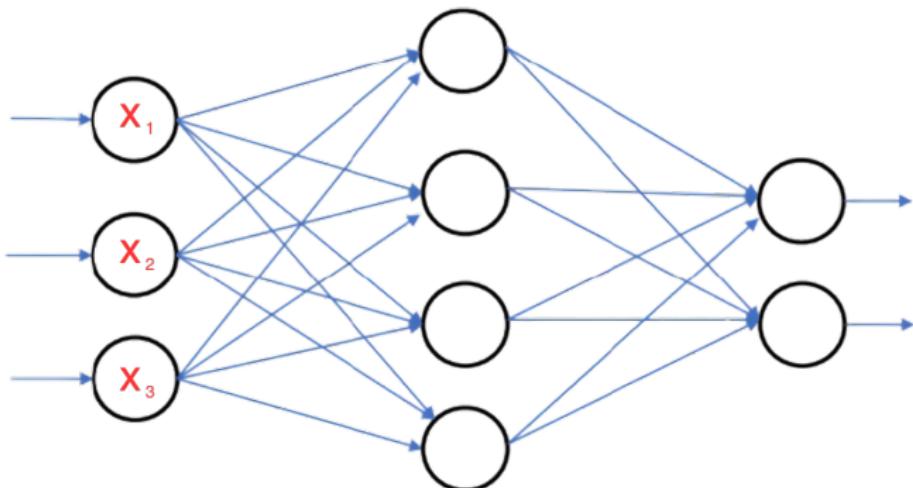
# Neural Networks



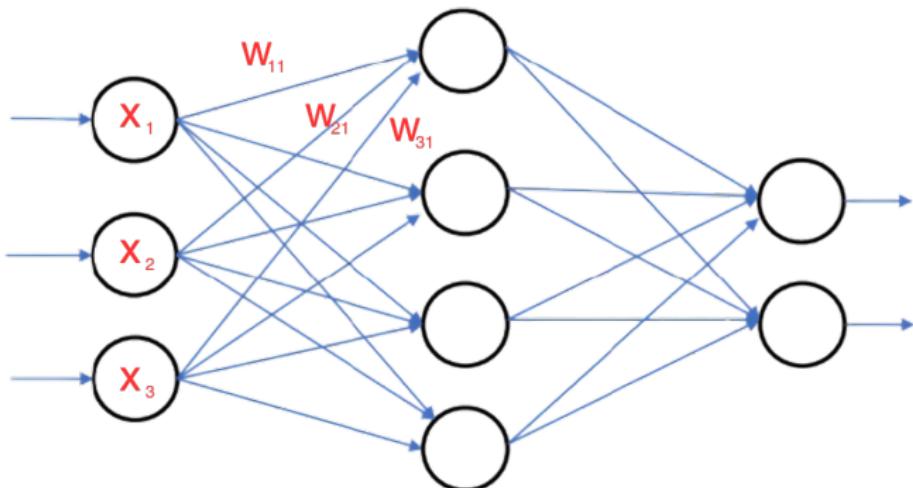
# Neural Networks



# Neural Networks

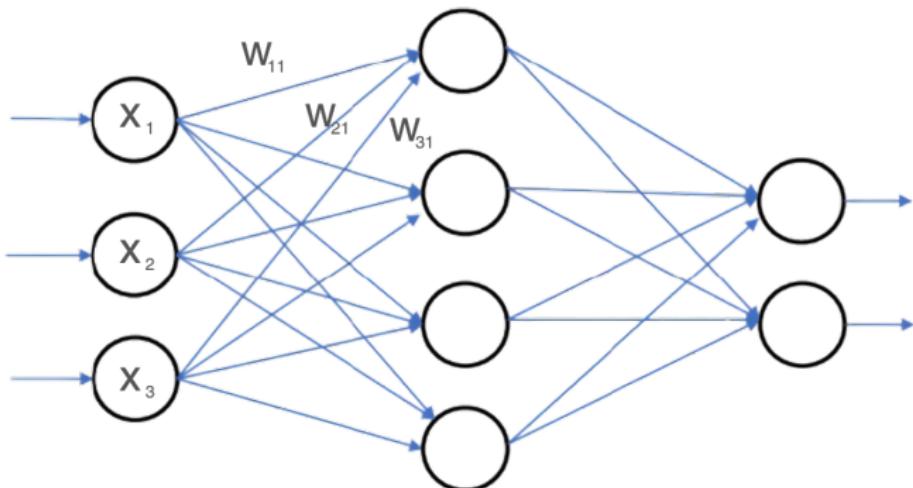


# Neural Networks



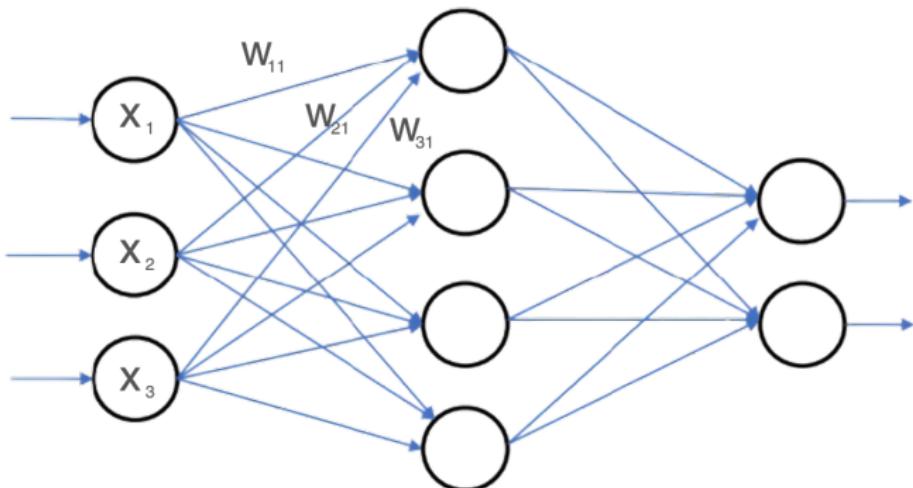
# Neural Networks

$$X_1 W_{11} + X_2 W_{21} + X_3 W_{31}$$



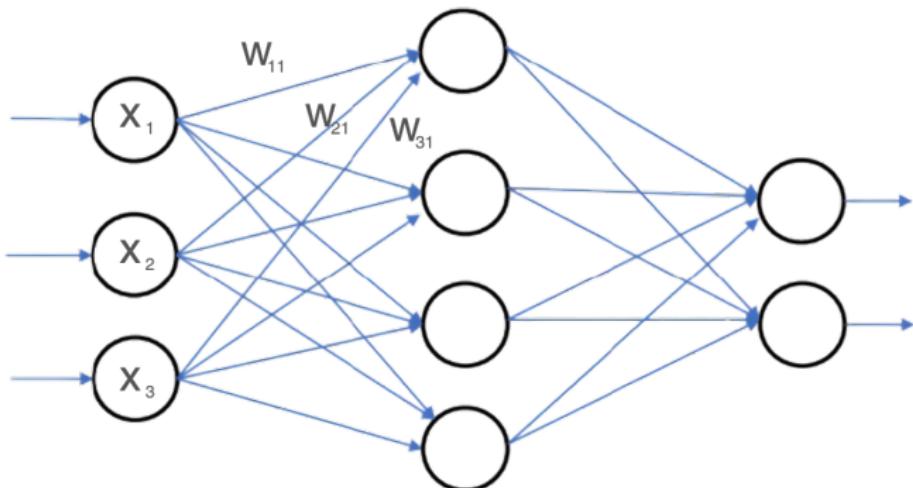
# Neural Networks

$$\sigma(x_1 w_{11} + x_2 w_{21} + x_3 w_{31})$$



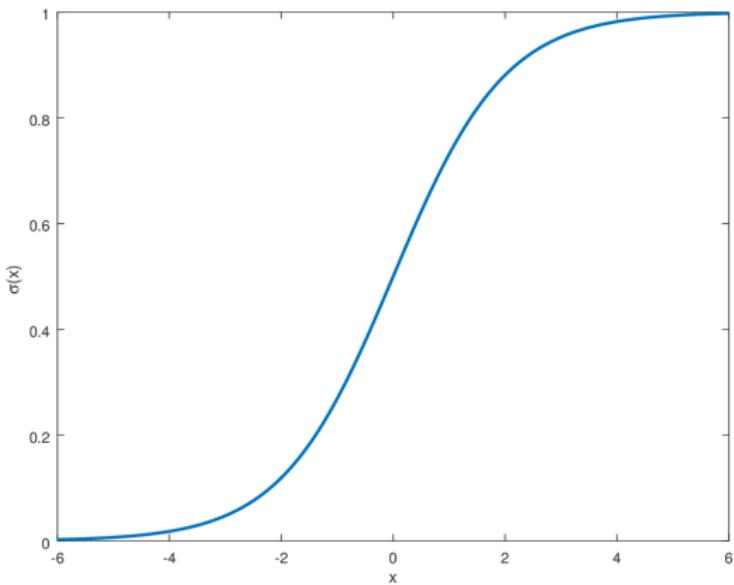
# Neural Networks

$$\sigma(x_1 w_{11} + x_2 w_{21} + x_3 w_{31}) = y_1$$

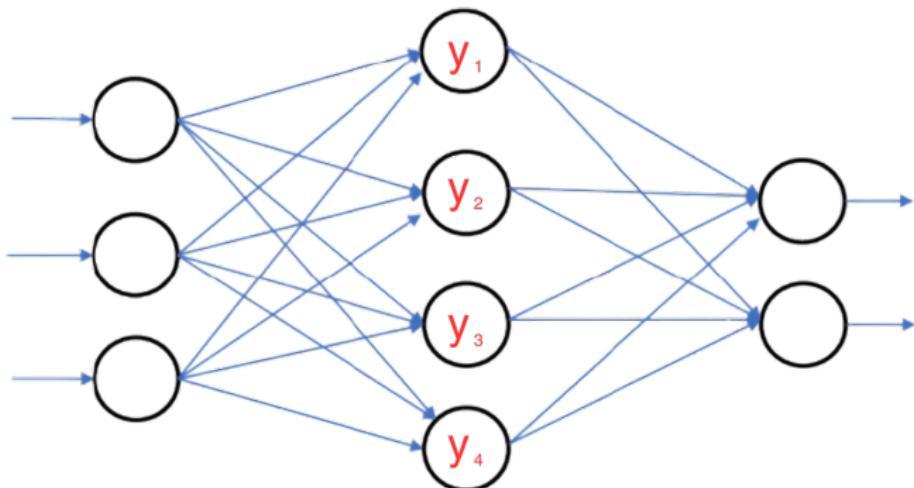


# Neural Networks-Sigmoid Function

$$\sigma(w_{11}x_1 + w_{21}x_2 + w_{31}x_3) = \frac{1}{1 + e^{-(w_{11}x_1 + w_{21}x_2 + w_{31}x_3)}}$$



# Neural Networks



## Neural Network Backpropagation

This algorithm computes gradients of the loss function with respect to each weight by applying the chain rule. For a neuron output  $a = \sigma(z)$ , where  $z = \sum_j w_j y_j + b$ , the gradient is:

$$\frac{\partial \mathcal{L}}{\partial w_i} = \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial y} \cdot \frac{\partial y}{\partial w_i}$$

Once the gradients are calculated, the weights are updated using:

$$w_i^{(t+1)} = w_i^{(t)} - \eta \frac{\partial \mathcal{L}}{\partial w_i}$$

where  $\eta$  is the learning rate.

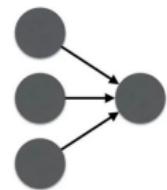
# GAN Example - Slanted Land



Slanted people

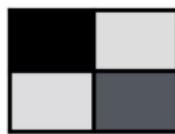


2x2 screens

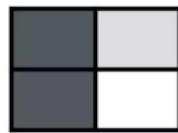


1-layer  
Neural networks

# Slanted Land - Faces



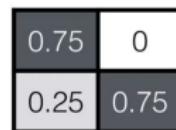
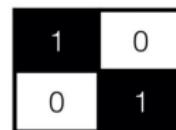
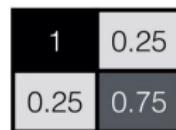
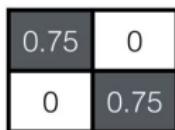
# Slanted Land - Noises



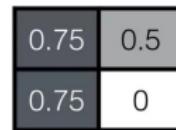
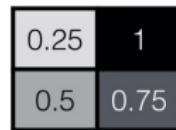
# Slanted Land - Faces and Noises



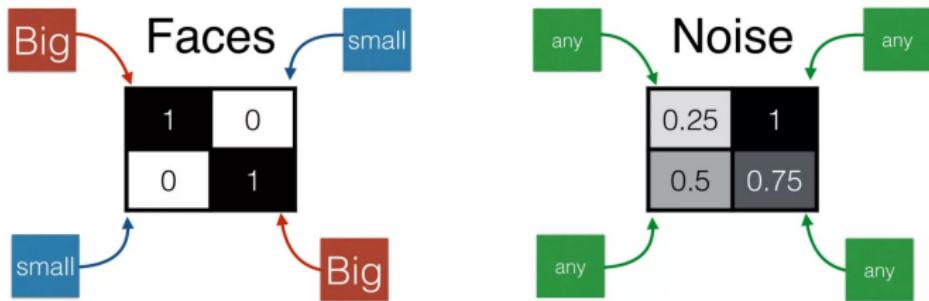
Faces



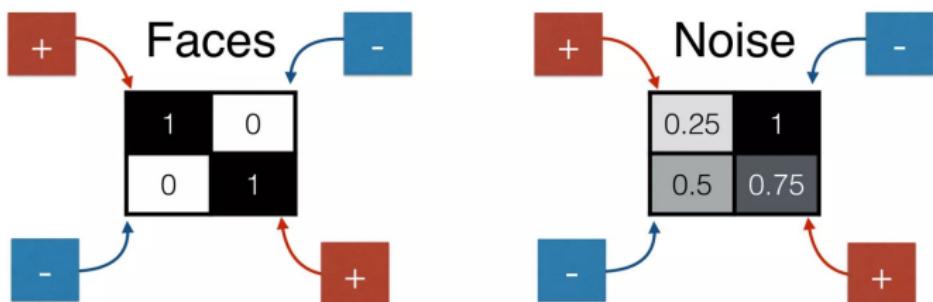
Noise



# Slanted Land - Faces and Noises



# Building Discriminator



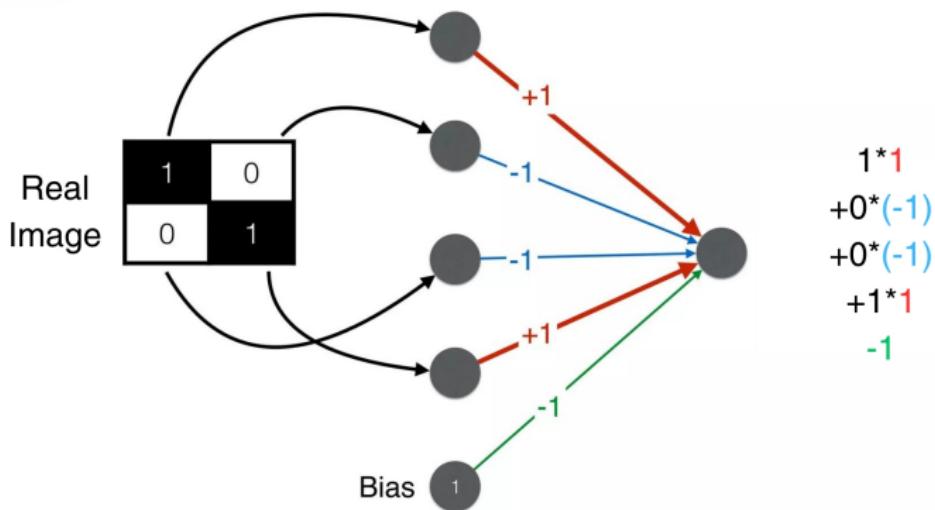
$$1*1 + 0*(-1) + 0*(-1) + 1*1 \\ = 2$$

Threshold = 1

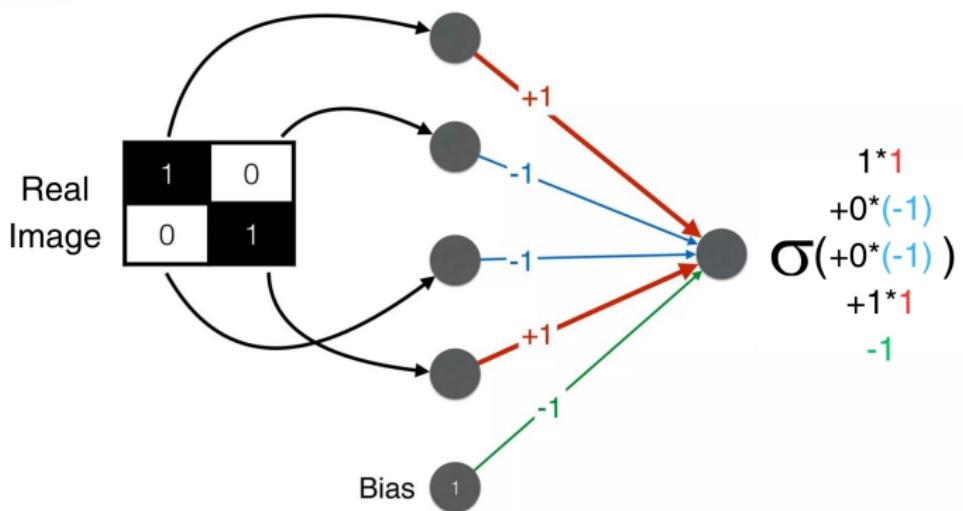
$$0.25*1 + 1*(-1) + 0.5*(-1) + 0.75*1 \\ = -0.5$$

More than 1: face. Less than 1: no face

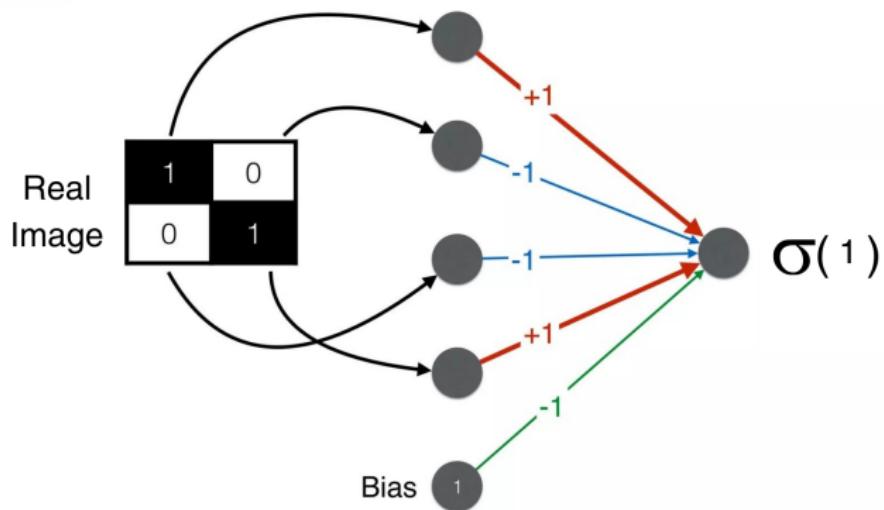
# Building Discriminator



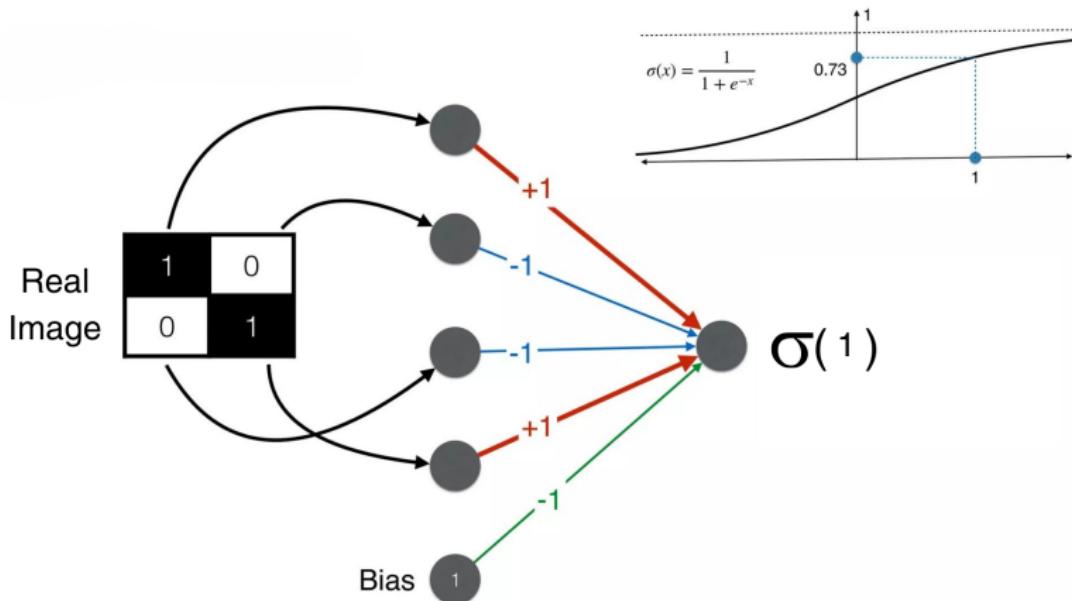
# Building Discriminator



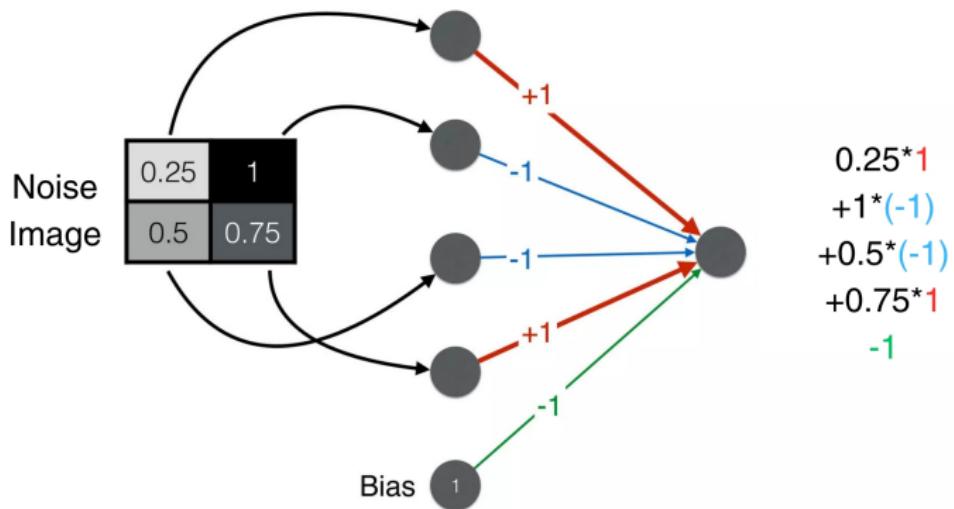
# Building Discriminator



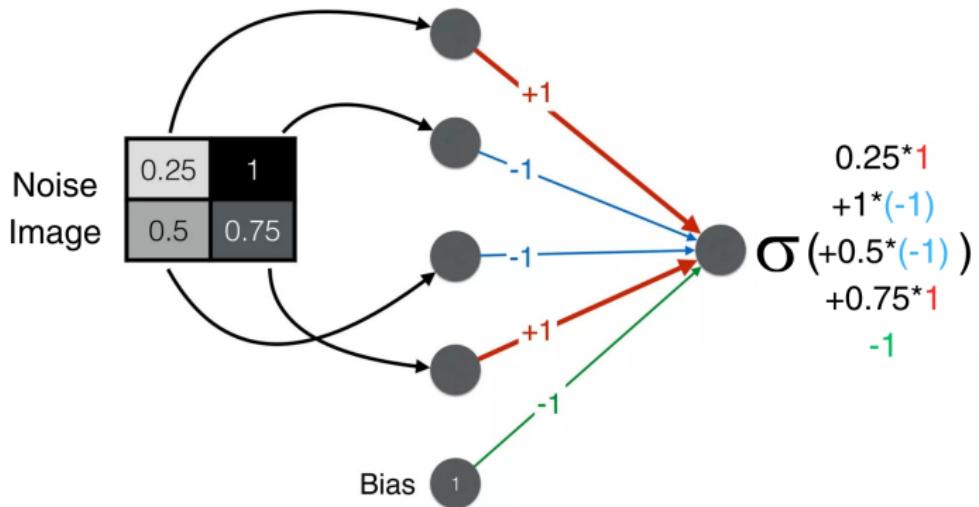
# Building Discriminator



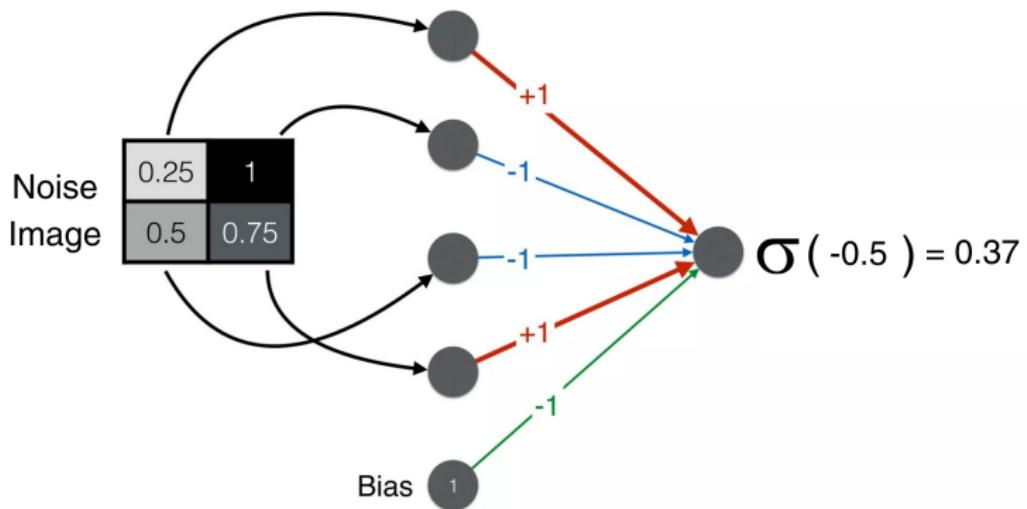
# Building Discriminator



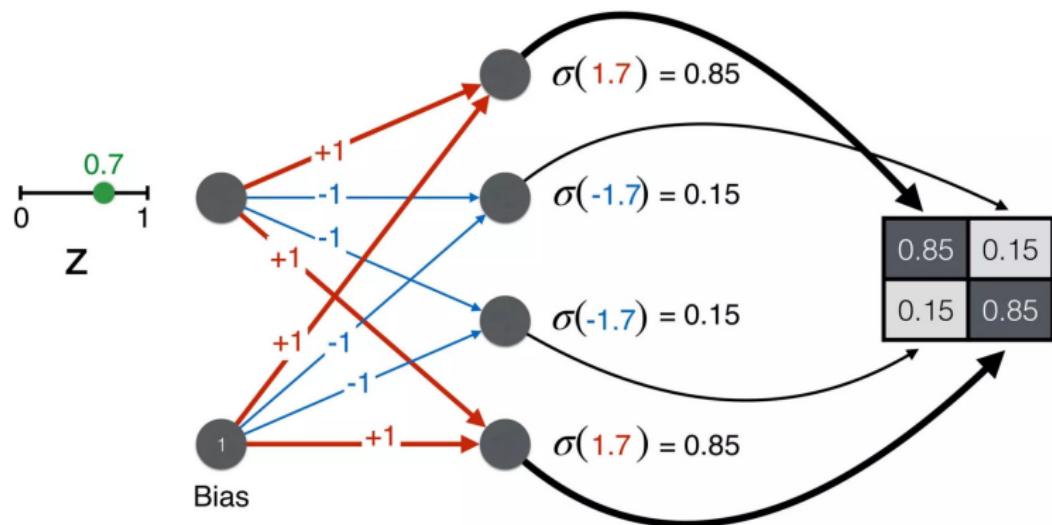
# Building Discriminator



# Building Discriminator



# Building Generator



# Log Loss Error

Label: 1

Prediction: 0.1

# Log Loss Error

Label: 1

Prediction: 0.1

Error: large

# Log Loss Error

Label: 1

Prediction: 0.1

Error: large

Label: 1

Prediction: 0.9

# Log Loss Error

$$\text{Error} = -\ln(\text{prediction})$$

Label: 1

Prediction: 0.1

Error: large

Label: 1

Prediction: 0.9

Error: small

# Log Loss Error

$$\text{Error} = -\ln(\text{prediction})$$

Label: 1

Prediction: 0.1      Error: large       $-\ln(0.1) = 2.3$

Label: 1

Prediction: 0.9      Error: small       $-\ln(0.9) = 0.1$

# Log Loss Error

Label: 0

Prediction: 0.1

# Log Loss Error

Label: 0

Prediction: 0.1

Error: small

# Log Loss Error

Label: 0

Prediction: 0.1

Error: small

Label: 0

Prediction: 0.9

# Log Loss Error

$$\text{Error} = -\ln(1 - \text{prediction})$$

Label: 0

Prediction: 0.1

Error: small

Label: 0

Prediction: 0.9

Error: large

# Log Loss Error

$$\text{Error} = -\ln(1 - \text{prediction})$$

Label: 0

Prediction: 0.1      Error: small       $-\ln(0.9) = 0.1$

Label: 0

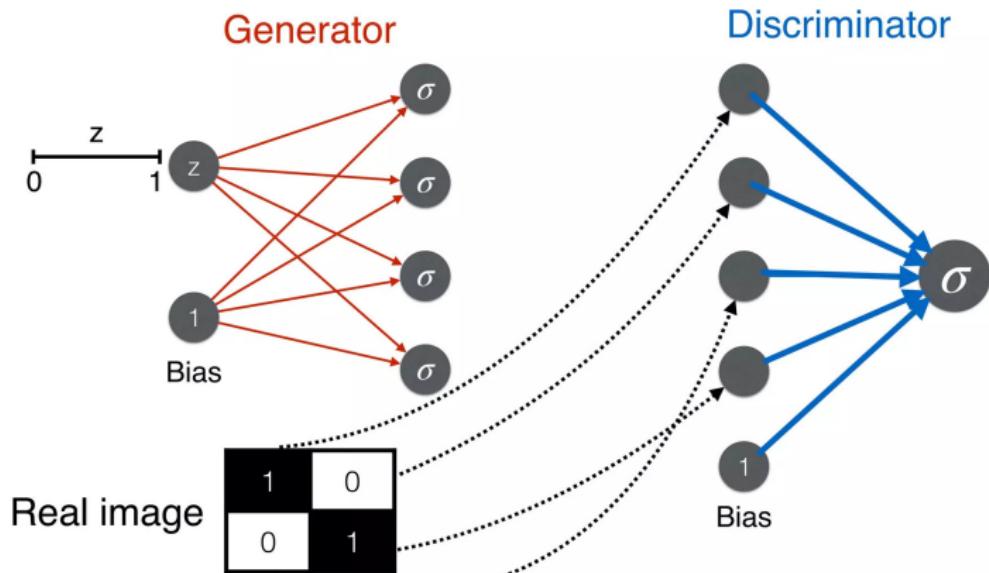
Prediction: 0.9      Error: large       $-\ln(0.1) = 2.3$

## GAN Objective Function

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))].$$

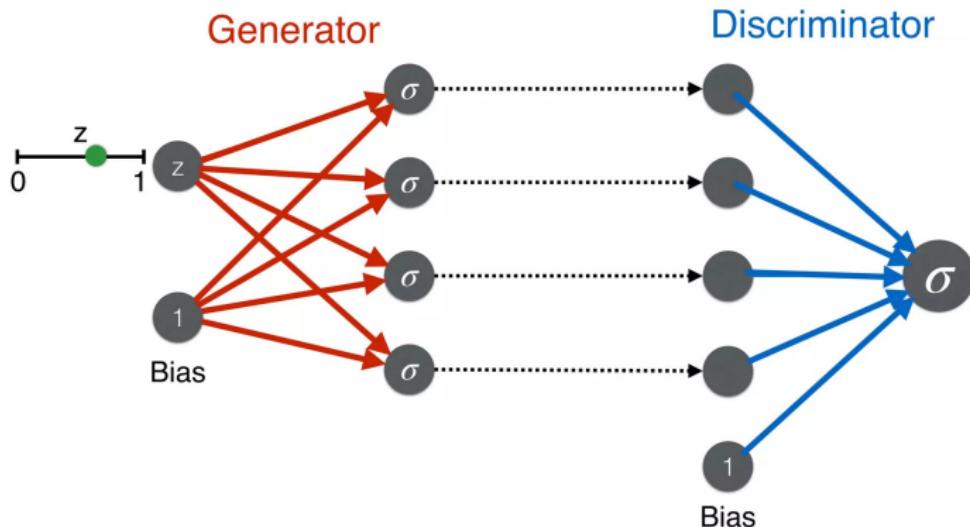
# Loss Function Optimization

$$\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})]$$



# Loss Function Optimization

$$\mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$



# Results



Figure 16: Original Image of Dwaipayan Haldar



Figure 17: Image Distorted by Manas Sarkar



Figure 18: Output Image of Dwaipayan Haldar restored by GAN

# Results



Figure 19: Original Ghibli Image of Dwaipayan Haldar



Figure 20: Ghibli Image Distorted by Pramit Khamrui



Figure 21: Output Ghibli Image of Dwaipayan Haldar restored by GAN

# Conclusion

# References

- Jireh Jam, Connah Kendrick, Kevin Walker, Vincent Drouard, Jison Gee-Sern Hsu, Moi Hoon Yap, A comprehensive review of past and present image inpainting methods, Computer Vision and Image Understanding, Volume 203, 2021.
- Efros, A.A., Leung, T.K., 1999. Texture synthesis by non-parametric sampling. In: Iccv. IEEE, p. 1033.
- Texture Synthesis by Non-parametric Sampling  
<https://people.eecs.berkeley.edu/~efros/research/NPS/alg.html>
- Thottam, I., 2015. The Cost of Conservation and Restoration. Art Business News,  
<http://artbusinessnews.com/2015/12/the-cost-of-conservation-and-restoration/>
- Liu, G., Reda, F.A., Shih, K.J., Wang, T.-C., Tao, A., Catanzaro, B., 2018a. Image inpainting for irregular holes using partial convolutions. arXiv preprint arXiv: 1804.07723.
- Serrano, L. (2019, November 11). Generative adversarial networks (GANs) [SlideShare presentation]. SlideShare.  
<https://www.slideshare.net/slideshow/generative-adversarial-networks-gans-238846644/238846644>
- Criminisi, A., Pérez, P., Toyama, K., 2004. Region filling and object removal by exemplar-based image inpainting. IEEE Trans. Image Process. 13 (9), 1200–1212.

# Thank You!

Questions?